

Bradley Voytek, Ph.D.
UC San Diego
Cognitive and Neural Dynamics Laboratory

Department of Cognitive Science
Neurosciences Graduate Program
Halicioğlu Data Science Institute

bvoytek@ucsd.edu
@bradleyvoytek

HALICIOĞLU DATA SCIENCE INSTITUTE






COGS 108

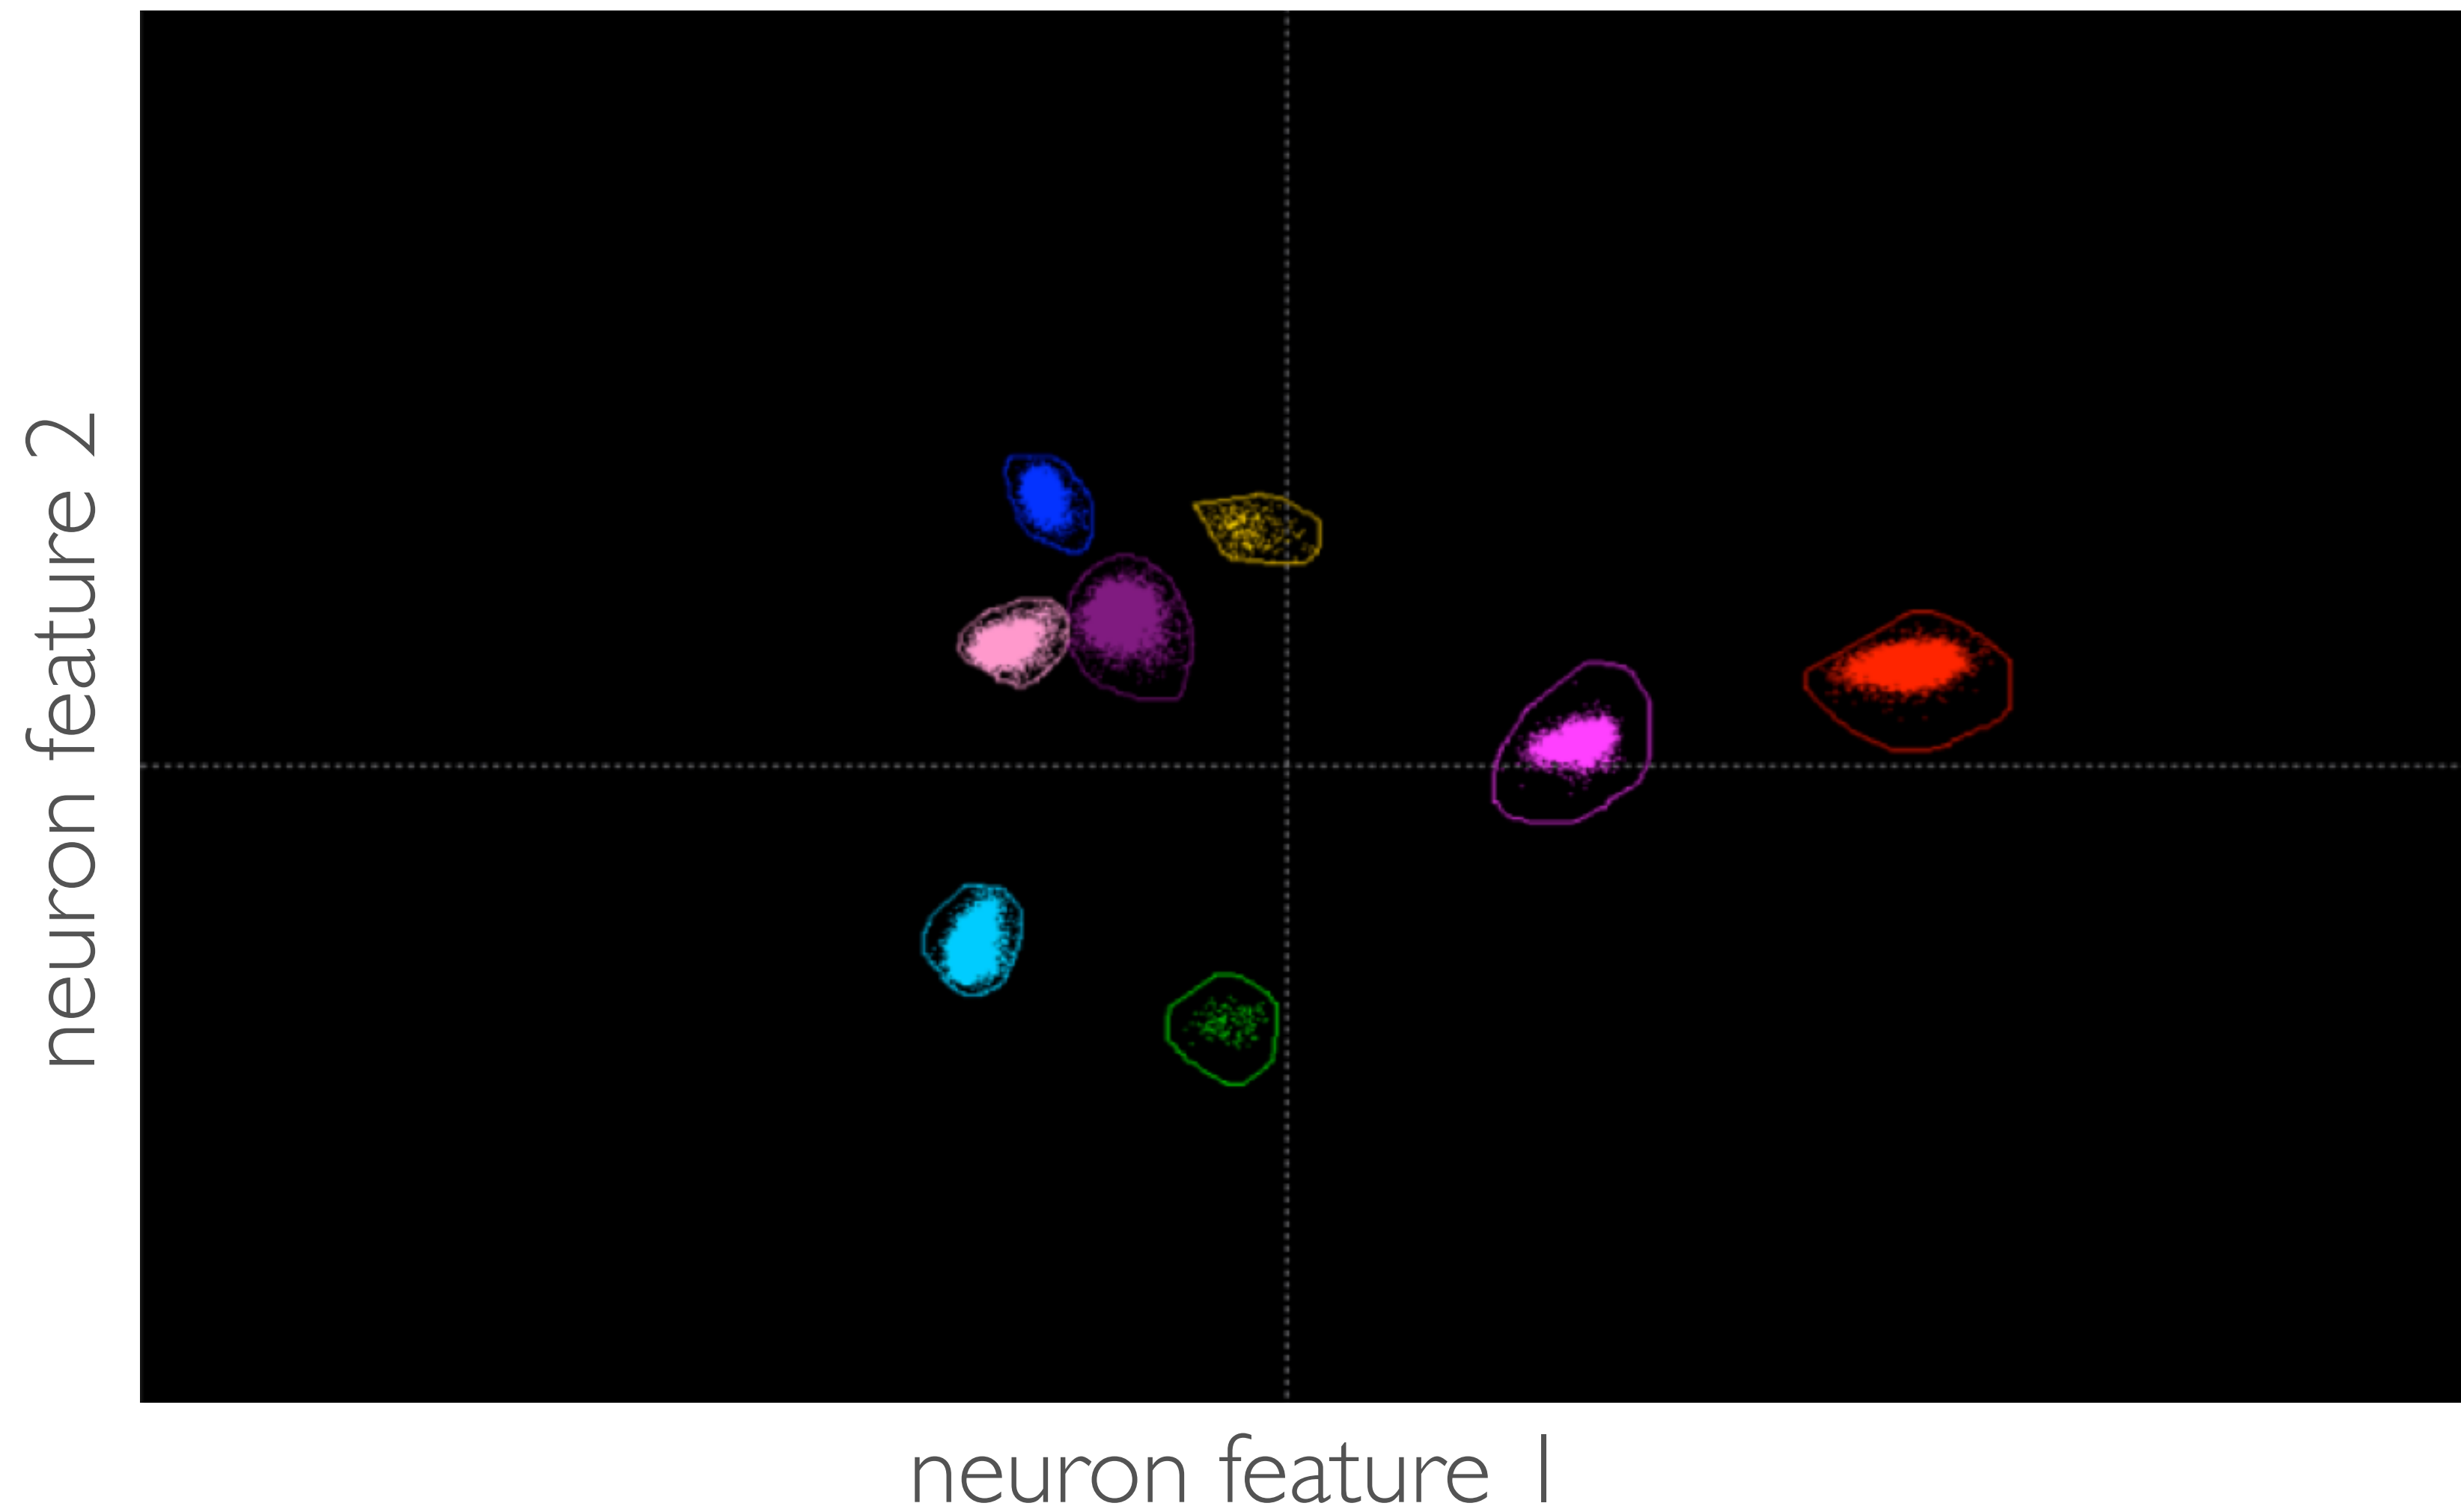
Data Science in Practice

Dimensionality reduction and clustering

Clustering



Clustering



Why clustering?

- Unlabeled data – unsupervised learning
 - If it is labeled, we'll use classification!
- Want to find groups!

Why clustering?

- Data Compression – We could want to reduce the dimensionality in the data.
 - Dimensionality – An aspect or a feature of something is a dimension.

How do we cluster?

- A lot of different ways we could, but we will start with **K-Means**.


k-means

- A simple but effective clustering algorithm
- Partitions the data in to K disjoint sets (clusters)
 - disjoint - no overlap
- We then run an iterative batch algorithm
 - batch – to do to all of the items (batch of cookies)


k-means

- Start – Initialize a guess of k centers/“means”
 - notation – j simply refers to which of the k centers.
- $S^{(j)}$ is all points closest to $\mu^{(j)}$
 - So each “mean” gets all of the points close to it in its set.
- Update the means:
$$\mu^{(j)} = \frac{1}{N_j} \sum_{n \in S^{(j)}} \mathbf{x}^{(n)}$$
- Continue until there is no change in the means or you decide to stop.


Initialize means




Initialize means




Label points




Connect adjacent means




Bisect




Label points




Move means to centroids




Move means to centroids




If means moved, repeat!




Bisect




Label points




Connect adjacent




Bisect




Label points




Connect adjacent




Bisect



Label points



Finished!



When do we reduce dimensionality?

Alicia is researching the resistance of plants in different soil to bugs. She has 10 each of 5 different types of plants. She has counted the number of bugs on each once per week for 10 weeks. Should Alicia analyze her bug count data with a clustering analysis? Why?

When do we reduce dimensionality?

Pat is studying cell types and morphologies in the basal forebrain of mice. They recorded the waveform amplitude, shapes, and firing patterns of 400 neurons. Should Pat run a clustering analysis?



When do we reduce dimensionality?

Linh is working at Apple. They have data on the accounts for all of their users, including what programs they use, how often, and for how long. They've been tasked with mining the data to find patterns that could affect marketing strategies. Should Linh run a clustering analysis?



COGS 108

Data Science in Practice

Support Vector Machines


SVM

**When classifying isn't easily linear,
*make it linear.***

SVM




SVM




SVM

These boats will be misclassified as houses




SVM




Shoreline is your
decision boundary

SVM




NOT LINEAR

SVM




SVM




Classify!

SVM



Classify!


SVM



Classify!

DONE


SVM



Classify!


~~DONE~~

SVM




This is not the
optimal separator

SVM




How do we find
the “optimal”?

SVM




These are your
“support vectors”

SVM




Support vectors are
the points nearest to
the plane

SVM




These are your
margins

SVM



Goal: maximize the margins (m) such that the decision boundary is as far away from the data of both classes as possible

SVM




Classify!

DONE

SVM


Okay... but...

SVM




Classify!

SVM




SVM



Classify!

SVM






Classify!

SVM




fy!

SVM




SVM




If we have 3 training points in 2D space, no matter how the points are labeled we can perfectly classify them using a linear algorithm

SVM




But add a fourth point, and this can fail

SVM




SVM: Total hack to turn data that are not optimally suited for linear classification into a linearly classifiable problem by warping them into higher dimensions

SVM




If we have N training points in $(N-1)D$ space, no matter how the points are labeled we can perfectly classify them using a linear algorithm

SVM



But rarely do you
need to go to $N-1$
space to do this

SVM



Classify!

Common SVM kernels

Examples:

$$K(\vec{x}_i, \vec{x}_j) = \vec{x}_i \cdot \vec{x}_j$$

Linear kernel

$$K(\vec{x}_i, \vec{x}_j) = \exp(-\gamma \|\vec{x}_i - \vec{x}_j\|^2)$$

Gaussian kernel

$$K(\vec{x}_i, \vec{x}_j) = \exp(-\gamma \|\vec{x}_i - \vec{x}_j\|)$$

Exponential kernel

$$K(\vec{x}_i, \vec{x}_j) = (p + \vec{x}_i \cdot \vec{x}_j)^q$$

Polynomial kernel

$$K(\vec{x}_i, \vec{x}_j) = (p + \vec{x}_i \cdot \vec{x}_j)^q \exp(-\gamma \|\vec{x}_i - \vec{x}_j\|^2)$$

Hybrid kernel

$$K(\vec{x}_i, \vec{x}_j) = \tanh(k\vec{x}_i \cdot \vec{x}_j - \delta)$$

Sigmoidal

COGS 108

Data Science in Practice

Geospatial Analyses

Geocoding

Geocoding is the process of converting addresses (like a street address) into geographic coordinates (like latitude and longitude), which you can use to place markers on a map, or position the map.

Google Maps API

Client-side geocoding, which is executed in the browser, generally in response to user action. The Google Maps JavaScript API provides classes that make the requests for you.

Google Maps API

When to use client-side geocoding

The short answer is "almost always." The reasons are:

- Client-side request and response provide a faster, more interactive experience for users.
- A client-side request can include information that improves geocoding quality: user language, region, and viewport.

Google Maps API

HTTP server-side geocoding, which allows your server to directly query Google's servers for geocodes. The Google Maps Geocoding API is the web service that provides this functionality. Typically, you integrate this service with other code that is running server-side.

Google Maps API

When to use server-side geocoding

Server-side geocoding is best used for applications that require you to geocode addresses without input from a client. A common example is when you get a dataset that comes independently of user input, for instance if you have a fixed, finite, and known set of addresses that need geocoding. Server-side geocoding can also be useful as a backup for when client-side geocoding fails.

Google Maps API

The Google Maps Geocoding API has the following limits in place:

Standard Usage Limits

Users of the standard API:

- 2,500 free requests per day, calculated as the sum of **client-side** and server-side queries.
- 50 requests per second, calculated as the sum of **client-side** and server-side queries.

Geocoding

Reverse geocoding is the process of converting geographic coordinates into a human-readable address. The Google Maps Geocoding API's reverse geocoding service also lets you find the address for a given place ID.

Geocoding

Jupyter Demo!

Administrative stuff

- **<http://cape.ucsd.edu/>**
- **COGS: <https://goo.gl/dmujYA>**
- The cognitive science department runs its own evaluations. You can submit evaluations for the instructor, your section TA, and/or any TA you had enough interaction with to provide feedback. Please do fill these out - they are very useful for us to get feedback on our teaching, and provide feedback to the department, and this is especially useful for working on this new course.

Bradley Voytek, Ph.D.
UC San Diego
Cognitive and Neural Dynamics Laboratory

Department of Cognitive Science
Neurosciences Graduate Program
Halicioğlu Data Science Institute

bvoytek@ucsd.edu
@bradleyvoytek